

Task 1 – Constraints

1.1 Original Constraints

1.1.1 EMP

```
SELECT OWNER, CONSTRAINT_NAME, TABLE_NAME, SEARCH_CONDITION,  
INDEX_NAME FROM USER_CONSTRAINTS WHERE TABLE_NAME = 'EMP';
```

OWNER	CONSTRAINT_NAME	TABLE_NAME	SEARCH_CONDITION	INDEX_NAME
S4518282	PK_EMPNO	EMP		PK_EMPNO

Elapsed: 00:00:00.34

1.1.2 DEPT

```
SELECT OWNER, CONSTRAINT_NAME, TABLE_NAME, SEARCH_CONDITION,  
INDEX_NAME FROM USER_CONSTRAINTS WHERE TABLE_NAME = 'DEPT';
```

OWNER	CONSTRAINT_NAME	TABLE_NAME	SEARCH_CONDITION	INDEX_NAME
S4518282	PK_DEPTNO	DEPT		PK_DEPTNO

Elapsed: 00:00:00.30

1.1.3 PURCHASE

```
SELECT OWNER, CONSTRAINT_NAME, TABLE_NAME, SEARCH_CONDITION,  
INDEX_NAME FROM USER_CONSTRAINTS WHERE TABLE_NAME = 'PURCHASE';
```

OWNER	CONSTRAINT_NAME	TABLE_NAME	SEARCH_CONDITION	INDEX_NAME
S4518282	PK_PURCHASENO	PURCHASE		PK_PURCHASENO

Elapsed: 00:00:00.28

1.1.4 CLIENT

```
SELECT OWNER, CONSTRAINT_NAME, TABLE_NAME, SEARCH_CONDITION,  
INDEX_NAME FROM USER_CONSTRAINTS WHERE TABLE_NAME = 'CLIENT';
```

OWNER	CONSTRAINT_NAME	TABLE_NAME	SEARCH_CONDITION	INDEX_NAME
S4518282	PK_CLIENTNO	CLIENT		PK_CLIENTNO

Elapsed: 00:00:00.29

1.2 Creation of missing constraints

1.2.1 EMP CONSTRAINTS

```
ALTER TABLE EMP ADD CONSTRAINT CK_ENAME  
CHECK (ENAME IS NOT NULL);
```

Table altered.

Elapsed: 00:00:00.05

```
ALTER TABLE EMP ADD CONSTRAINT  
FK_DEPTNO FOREIGN KEY (DEPTNO)  
REFERENCES DEPT (DEPTNO);
```

Table altered.

Elapsed: 00:00:00.01

1.2.2 DEPT CONSTRAINTS

ALTER TABLE DEPT ADD CONSTRAINT
UN_DNAME UNIQUE (DNAME);

```
Table altered.  
Elapsed: 00:00:00.14
```

ALTER TABLE DEPT ADD CONSTRAINT
CK_DNAME CHECK (DNAME IS NOT NULL);

```
Table altered.  
Elapsed: 00:00:00.01
```

1.2.3 PURCHASE CONSTRAINTS

ALTER TABLE PURCHASE ADD CONSTRAINT
CK_RECEIPTNO CHECK (RECEIPTNO IS NOT NULL);

```
Table altered.  
Elapsed: 00:00:00.02
```

ALTER TABLE PURCHASE ADD CONSTRAINT
CK_AMOUNT CHECK (AMOUNT IS NOT NULL);

```
Table altered.  
Elapsed: 00:00:00.02
```

ALTER TABLE PURCHASE ADD CONSTRAINT
CK_SERVICETYPE CHECK (SERVICETYPE IN
('Training', 'Data Recovery', 'Consultation', 'Software
Installation', 'Software Repair'));

```
Table altered.  
Elapsed: 00:00:00.01
```

ALTER TABLE PURCHASE ADD CONSTRAINT
CK_PAYMENTTYPE CHECK (PAYMENTTPYE IN
('Debit', 'Cash', 'Credit'));

```
Table altered.  
Elapsed: 00:00:00.02
```

ALTER TABLE PURCHASE ADD CONSTRAINT
CK_GST CHECK (GST IN ('Yes', 'No'));

```
Table altered.  
Elapsed: 00:00:00.01
```

ALTER TABLE PURCHASE ADD CONSTRAINT
FK_EMPNO FOREIGN KEY (SERVEDBY)
REFERENCES EMP (EMPNO);

```
Table altered.  
Elapsed: 00:00:00.01
```

ALTER TABLE PURCHASE ADD CONSTRAINT
FK_CLIENTNO FOREIGN KEY (CLIENTNO)
REFERENCES CLIENT (CLIENTNO);

```
Table altered.  
Elapsed: 00:00:00.01
```

1.2.4 CLIENT CONSTRAINTS

ALTER TABLE CLIENT ADD CONSTRAINT
CK_CNAME CHECK (CNAME IS NOT NULL);

```
Table altered.  
Elapsed: 00:00:00.01
```

1.2.5 EMP

```
SELECT OWNER, CONSTRAINT_NAME, TABLE_NAME, SEARCH_CONDITION,
INDEX_NAME FROM USER_CONSTRAINTS WHERE TABLE_NAME = 'EMP';
```

OWNER	CONSTRAINT_NAME	TABLE_NAME	SEARCH_CONDITION	INDEX_NAME
S4518282	FK_DEPTNO	EMP	ENAME IS NOT NULL	PK_EMPNO
S4518282	CK_ENAME	EMP		
S4518282	PK_EMPNO	EMP		

Elapsed: 00:00:00.00

1.2.6 DEPT

```
SELECT OWNER, CONSTRAINT_NAME, TABLE_NAME, SEARCH_CONDITION,
INDEX_NAME FROM USER_CONSTRAINTS WHERE TABLE_NAME = 'DEPT';
```

OWNER	CONSTRAINT_NAME	TABLE_NAME	SEARCH_CONDITION	INDEX_NAME
S4518282	CK_DNAME	DEPT	DNAME IS NOT NULL	UN_DNAME PK_DEPTNO
S4518282	UN_DNAME	DEPT		
S4518282	PK_DEPTNO	DEPT		

Elapsed: 00:00:00.00

1.2.7 PURCHASE

```
SELECT OWNER, CONSTRAINT_NAME, TABLE_NAME, SEARCH_CONDITION,
INDEX_NAME FROM USER_CONSTRAINTS WHERE TABLE_NAME = 'PURCHASE';
```

OWNER	CONSTRAINT_NAME	TABLE_NAME	SEARCH_CONDITION	INDEX_NAME
S4518282	FK_EMPNO	PURCHASE	RECEIPTNO IS NOT NULL	
S4518282	FK_CLIENTNO	PURCHASE		
S4518282	CK_RECEIPTNO	PURCHASE		
S4518282	CK_AMOUNT	PURCHASE	AMOUNT IS NOT NULL SERVICETYPE IN ('Training', 'Data Recovery', 'Consultation', 'Software Installation')	
S4518282	CK_SERVICETYPE	PURCHASE		
S4518282	CK_PAYMENTTYPE	PURCHASE	PAYMENTTYPE IN ('Debit', 'Cash', 'Credit')	
S4518282	CK_GST	PURCHASE	GST IN ('Yes', 'No')	PK_PURCHASENO
S4518282	PK_PURCHASENO	PURCHASE		

8 rows selected.

Elapsed: 00:00:00.01

1.2.8 CLIENT

```
SELECT OWNER, CONSTRAINT_NAME, TABLE_NAME, SEARCH_CONDITION,
INDEX_NAME FROM USER_CONSTRAINTS WHERE TABLE_NAME = 'CLIENT';
```

OWNER	CONSTRAINT_NAME	TABLE_NAME	SEARCH_CONDITION	INDEX_NAME
S4518282	CK_CNAME	CLIENT	CNAME IS NOT NULL	PK_CLIENTNO
S4518282	PK_CLIENTNO	CLIENT		

Elapsed: 00:00:00.00

Task 2 – *Triggers*

2.1 Top Client

```
SELECT C.CLIENTNO, CNAME, SUM(AMOUNT) FROM CLIENT C, PURCHASE P
WHERE C.CLIENTNO = P.CLIENTNO
GROUP BY C.CLIENTNO, CNAME
HAVING SUM(AMOUNT) >= ALL(
    SELECT SUM(AMOUNT) FROM CLIENT C, PURCHASE P
    WHERE C.CLIENTNO = P.CLIENTNO
    GROUP BY C.CLIENTNO, CNAME);
```

CLIENTNO	CNAME	SUM(AMOUNT)
24535	Sally Moon	20100

Elapsed: 00:00:03.37

2.2 TOP_DISCOUNT Trigger

If the tuple that is being inserted into the PURCHASE table has the CLIENTNO = 24535 (Sally Moon as she is the top client from task 2.1) then we must make the AMOUNT be 85% of what it's supposed to be.

```
CREATE OR REPLACE TRIGGER TOP_DISCOUNT
BEFORE INSERT ON "PURCHASE"
FOR EACH ROW
BEGIN
    IF (:NEW.CLIENTNO = 24535)
    THEN :NEW.AMOUNT := 0.85* :NEW.AMOUNT;
    END IF;
END;
/
```

Trigger created.
Elapsed: 00:00:00.02

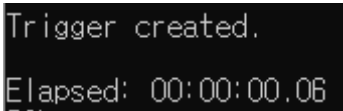
2.3 SUNCHINE_DEPT Trigger

Firstly, put all employee working in the department 'SALES – Sunshine' into a temporary table vEMPNO and check if the tuple that is being inserted into PURCHASE. If the client is being SERVEDBY an employee within vEMPNO then we must change the PAYMENTTYPE = 'Cash'. Also if the SERVICETYPE = 'Data Recovery' we must change the AMOUNT to 70% of what it's supposed to be.

```
CREATE OR REPLACE TRIGGER TOP_DISCOUNT
  BEFORE INSERT ON "PURCHASE"
  FOR EACH ROW
  DECLARE
    vEMPNO EMP.EMPNO%TYPE;
BEGIN
  SELECT EMPNO INTO vEMPNO FROM EMP WHERE DEPTNO = (
    SELECT DEPTNO FROM DEPT WHERE DNAME = 'SALES – Sunshine');

  IF (:NEW.SERVEDBY IN (vEMPNO))
  THEN :NEW.PAYMENTTYPE := 'Cash';
  END IF;

  IF (:NEW.SERVICETYPE = 'Data Recovery')
  THEN :NEW.AMOUNT := 0.70* :NEW.AMOUNT;
  END IF;
END;
/
```



```
Trigger created.
Elapsed: 00:00:00.06
```

Task 3 – Views

3.1 V_DEPT_AMOUNT

```
CREATE VIEW V_DEPT_AMOUNT AS
SELECT D.DEPTNO, D.DNAME, MAX(AMOUNT) MAXA, MIN(AMOUNT) MINA,
AVG(AMOUNT) AVGA, SUM(AMOUNT) SUMA
FROM DEPT D, PURCHASE P, EMP E
WHERE P.SERVEDBY = E.EMPNO AND D.DEPTNO = E.DEPTNO
GROUP BY D.DEPTNO, D.DNAME;
```

```
View created.
Elapsed: 00:00:00.01
```

3.2 MV_DEPT_AMOUNT

```
CREATE MATERIALIZED VIEW V_DEPT_AMOUNT AS BUILD IMMEDIATE AS
SELECT D.DEPTNO, D.DNAME, MAX(AMOUNT) MAXA, MIN(AMOUNT) MINA,
AVG(AMOUNT) AVGA, SUM(AMOUNT) SUMA
FROM DEPT D, PURCHASE P, EMP E
WHERE P.SERVEDBY = E.EMPNO AND D.DEPTNO = E.DEPTNO
GROUP BY D.DEPTNO, D.DNAME;
```

```
Materialized view created.
Elapsed: 00:00:00.41
```

3.3 VIEW EXECUTION AND EXPLANATION

```
SELECT * FROM V_DEPT_AMOUNT;
```

DEPTNO	DNAME	MAXA	MINA	AVGA	SUMA
30	SALES - Sunflower	1000	50	528.336968	968970
40	SALES - Hercules	1000	50	535.761089	1062950
50	SALES - Neptune	1000	50	517.578053	1674365
20	SALES - Sunshine	1000	50	522.126719	1063050
10	SALES - Glorious	1000	50	522.730769	475685

Elapsed: 00:00:00.05

```
SELECT * FROM MV_DEPT_AMOUNT;
```

DEPTNO	DNAME	MAXA	MINA	AVGA	SUMA
30	SALES - Sunflower	1000	50	528.336968	968970
40	SALES - Hercules	1000	50	535.761089	1062950
50	SALES - Neptune	1000	50	517.578053	1674365
20	SALES - Sunshine	1000	50	522.126719	1063050
10	SALES - Glorious	1000	50	522.730769	475685

Elapsed: 00:00:00.01

As shown above in the execution of both queries which retrieve the data from the virtual or materialized view, the query execution time of the materialized view is faster than the virtual view. Although small it is still noticeable (00.05 vs 00.01 for materialized and virtual views respectively).

Query Execution Plan:

```
EXPLAIN PLAN FOR SELECT * FROM
V_DEPT_AMOUNT;
```

```
Explained.
Elapsed: 00:00:00.03
```

```
EXPLAIN PLAN FOR SELECT * FROM
MV_DEPT_AMOUNT;
```

```
Explained.
Elapsed: 00:00:00.00
```

```
SELECT PLAN_TABLE_OUTPUT FROM TABLE (DBMS_XPLAN.DISPLAY); -- VIRTUAL VIEW
```

```
PLAN_TABLE_OUTPUT
Plan hash value: 3964501273
```

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		18	630	30 (10)	00:00:01
1	HASH GROUP BY		18	630	30 (10)	00:00:01
* 2	HASH JOIN		9999	341K	28 (4)	00:00:01
3	MERGE JOIN		75	2025	6 (17)	00:00:01
4	TABLE ACCESS BY INDEX ROWID	DEPT	5	100	2 (0)	00:00:01
5	INDEX FULL SCAN	PK_DEPTNO	5		1 (0)	00:00:01

```
SELECT PLAN_TABLE_OUTPUT FROM TABLE (DBMS_XPLAN.DISPLAY); -- MATERIALIZED VIEW
```

```
PLAN_TABLE_OUTPUT
Plan hash value: 3751959186
```

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		5	270	3 (0)	00:00:01
1	MAT_VIEW ACCESS FULL	MV_DEPT_AMOUNT	5	270	3 (0)	00:00:01

```
8 rows selected.
Elapsed: 00:00:00.06
```

The reason can be seen in the execution plans above. As virtual views are not physically stored within the database the views/queries must be recomputed each time the query is ran. Thus oracle has to perform query optimization by choosing to use efficient algorithms such as 'HASH JOIN', 'MERGE JOIN' & Taking use of the index on the primary key (DEPTNO). However, because materialized views are stored physically oracle has gone ahead and accessed the view/table directly, which can be seen by the 'MAT_VIEW ACCESS FULL', thus no query modification is performed which will overall improve the speed of query processing. Ultimately, the 'Time' column in the execution plan immediately shows the query processing time differences.

3.4 DEPT_EMP_AMOUNT

3.4.1 V_DEPT_EMP_AMOUNT

CREATE VIEW V_DEPT_EMP_AMOUNT **AS**

SELECT D.DEPTNO, EMPNO, COUNT(*) CA, AVG(AMOUNT) AVGA, MAX(AMOUNT) MAXA, SUM(AMOUNT) SUMA

FROM DEPT D, EMP E, PURCHASE P

WHERE D.DEPTNO = E.DEPTNO AND P.SERVEDBY = E.EMPNO

GROUP BY D.DEPTNO, EMPNO

ORDER BY D.DEPTNO ASC, SUMA DESC;

```
View created.  
Elapsed: 00:00:00.01
```

3.4.2 MV_DEPT_EMP_AMOUNT

CREATE MATERIALIZED VIEW MV_DEPT_EMP_AMOUNT **BUILD IMMEDIATE AS**

SELECT D.DEPTNO, EMPNO, COUNT(*) CA, AVG(AMOUNT) AVGA, MAX(AMOUNT) MAXA, SUM(AMOUNT) SUMA

FROM DEPT D, EMP E, PURCHASE P

WHERE D.DEPTNO = E.DEPTNO AND P.SERVEDBY = E.EMPNO

GROUP BY D.DEPTNO, EMPNO

ORDER BY D.DEPTNO ASC, SUMA DESC;

```
Materialized view created.  
Elapsed: 00:00:00.08
```

3.5 VIEW EXECUTION AND EXPLANATION

SELECT * FROM V_DEPT_EMP_AMOUNT;

DEPTNO	EMPNO	CA	AVGA	MAXA	SUMA
10	1065	166	543.614458	990	90240
10	1007	167	502.664671	990	83945
10	1015	150	544.233333	995	81635
10	1031	144	524.305556	1000	75500
10	1009	138	546.195652	1000	75375
10	1055	144	472.636889	985	68060
10	1071	1	930	930	930
20	1022	167	533.952096	995	89170
20	1049	155	550.548387	1000	85335
20	1039	158	529.683544	995	83690
50	1023	138	491.594203	995	67840
50	1046	136	480.735294	980	65380
50	1050	125	509.08	995	63635
50	1030	121	518.636364	995	62755
50	1016	126	477.777778	970	60200
50	1025	116	500.775862	990	58090
50	1070	116	476.724138	995	55300
50	1075	1	450	450	450

75 rows selected.
Elapsed: 00:00:00.07

SELECT * FROM MV_DEPT_EMP_AMOUNT;

DEPTNO	EMPNO	CA	AVGA	MAXA	SUMA
10	1065	166	543.614458	990	90240
10	1007	167	502.664671	990	83945
10	1015	150	544.233333	995	81635
10	1031	144	524.305556	1000	75500
10	1009	138	546.195652	1000	75375
10	1055	144	472.638889	985	68060
10	1071	1	930	930	930
20	1022	167	533.952096	995	89170
20	1049	155	550.548387	1000	85335
20	1039	158	529.683544	995	83690
50	1023	138	491.594203	995	67840
50	1046	136	480.735294	980	65380
50	1050	125	509.08	995	63635
50	1030	121	518.636364	995	62755
50	1016	126	477.777778	970	60200
50	1025	116	500.775862	990	58090
50	1070	116	476.724138	995	55300
50	1075	1	450	450	450

75 rows selected.
Elapsed: 00:00:00.04

Query Execution Plan:

EXPLAIN PLAN FOR SELECT * FROM
V_DEPT_EMP_AMOUNT;

Explained.
Elapsed: 00:00:00.01

EXPLAIN PLAN FOR SELECT * FROM
MV_DEPT_EMP_AMOUNT;

Explained.
Elapsed: 00:00:00.00

SELECT PLAN_TABLE_OUTPUT FROM TABLE (DBMS_XPLAN.DISPLAY); -- VIRTUAL VIEW

PLAN_TABLE_OUTPUT						
Plan hash value: 862709908						
Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		75	5850	27 (12)	00:00:01
1	VIEW	V_DEPT_EMP_AMOUNT	75	5850	27 (12)	00:00:01
2	SORT ORDER BY		75	4725	27 (12)	00:00:01
3	MERGE JOIN		75	4725	26 (8)	00:00:01
* 4	TABLE ACCESS BY INDEX ROWID	EMP	75	525	2 (0)	00:00:01
5	INDEX FULL SCAN	PK_EMPNO	75		1 (0)	00:00:01
PLAN_TABLE_OUTPUT						
* 6	SORT JOIN		75	4200	24 (9)	00:00:01
7	VIEW	VW_GBC_6	75	4200	23 (5)	00:00:01
8	HASH GROUP BY		75	600	23 (5)	00:00:01
9	TABLE ACCESS FULL	PURCHASE	9999	79992	22 (0)	00:00:01

SELECT PLAN_TABLE_OUTPUT FROM TABLE (DBMS_XPLAN.DISPLAY); -- MATERIALIZED VIEW

PLAN_TABLE_OUTPUT						
Plan hash value: 1444910117						
Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		75	3000	3 (0)	00:00:01
1	MAT_VIEW ACCESS FULL	MV_DEPT_EMP_AMOUNT	75	3000	3 (0)	00:00:01

8 rows selected.
Elapsed: 00:00:00.03

Similarly to task 3.1~3.3 the query execution time of the materialized view was faster than the virtual view (00.04 vs 00.07 respectively). This is again expected as materialized views are stored physically while the virtual views are not. Thus as shown in the execution plan for virtual view above, oracle uses query optimization for the query processing of the virtual view by using efficient algorithms such as 'HASH JOIN', 'MERGE JOIN' & even taking advantage of the index on the primary key (EMPNO). However when executing the query for the materialized view, oracle just needs to access the table stored physically thus making the execution of materialized views much faster.

Task 4 – Indexes

4.1 Receipt book

```
SELECT COUNT(*) FROM PURCHASE
WHERE SUBSTR(RECEIPTNO, 0, 3) IN (
    SELECT SUBSTR(RECEIPTNO, 0, 3) FROM PURCHASE
    GROUP BY SUBSTR(RECEIPTNO, 0, 3)
    HAVING COUNT(*) >= 10);
```

```
COUNT(*)
-----
7896
Elapsed: 00:00:00.04
```

4.2 Receipt book query execution plan/time and creation of index

```
EXPLAIN PLAN FOR SELECT COUNT(*) FROM PURCHASE
WHERE SUBSTR(RECEIPTNO, 0, 3) IN (
    SELECT SUBSTR(RECEIPTNO, 0, 3) FROM PURCHASE
    GROUP BY SUBSTR(RECEIPTNO, 0, 3)
    HAVING COUNT(*) >= 10);
```

```
Explained.
Elapsed: 00:00:00.03
```

SELECT PLAN_TABLE_OUTPUT FROM TABLE (DBMS_XPLAN.DISPLAY); --Before Index

```
PLAN_TABLE_OUTPUT
-----
Plan hash value: 1341602385

| Id | Operation | Name | Rows | Bytes | Cost (%CPU)| Time |
|----|-----|-----|-----|-----|-----|-----|
| 0 | SELECT STATEMENT | | 1 | 13 | 46 (5) | 00:00:01 |
| 1 | SORT AGGREGATE | | 1 | 13 | | |
| * 2 | HASH JOIN RIGHT SEMI | | 100 | 1300 | 46 (5) | 00:00:01 |
| 3 | VIEW | VW_NSO_1 | 497 | 3976 | 23 (5) | 00:00:01 |
| * 4 | FILTER | | | | | |
| 5 | HASH GROUP BY | | 497 | 2485 | 23 (5) | 00:00:01 |

PLAN_TABLE_OUTPUT
-----
| 6 | TABLE ACCESS FULL | PURCHASE | 9999 | 49995 | 22 (0) | 00:00:01 |
| 7 | TABLE ACCESS FULL | PURCHASE | 9999 | 49995 | 22 (0) | 00:00:01 |

Predicate Information (identified by operation id):
-----
 2 - access("SUBSTR(RECEIPTNO,0,3)"=SUBSTR(TO_CHAR("RECEIPTNO"),0,3))
 4 - filter(COUNT(*)>=10)

20 rows selected.
Elapsed: 00:00:00.03
```

```
CREATE INDEX BOOK_INDEX ON PURCHASE SUBSTR(RECEIPTNO, 0, 3);
```

```
Index created.
Elapsed: 00:00:00.06
```

Result of Query After index creation.

```
COUNT(*)
-----
7896
Elapsed: 00:00:00.01
```

```
SELECT PLAN_TABLE_OUTPUT FROM TABLE (DBMS_XPLAN.DISPLAY); --After Index
```

```
PLAN_TABLE_OUTPUT
-----
Plan hash value: 2292030079

| Id | Operation                      | Name      | Rows  | Bytes | Cost (%CPU)| Time     |
|----|-----|-----|-----|-----|-----|-----|-----|
| 0  | SELECT STATEMENT                |           |      1 |    13 |    22 (10) | 00:00:01 |
| 1  |   SORT AGGREGATE                |           |      1 |    13 |           |          |
|* 2  |    HASH JOIN RIGHT SEMI         |           |     100 |   1300 |    22 (10) | 00:00:01 |
| 3  |      VIEW                       | VW_NSO_1  |     497 |   3976 |    11 (10) | 00:00:01 |
|* 4  |        FILTER                   |           |      497 |   2485 |    11 (10) | 00:00:01 |
| 5  |          HASH GROUP BY          |           |      497 |   2485 |           |          |

PLAN_TABLE_OUTPUT
-----
| 6  |      INDEX FAST FULL SCAN       | BOOK_INDEX |    9999 |   49995 |    10 (0) | 00:00:01 |
| 7  |      INDEX FAST FULL SCAN       | BOOK_INDEX |    9999 |   49995 |    10 (0) | 00:00:01 |

Predicate Information (identified by operation id):
-----
   2 - access("SUBSTR(RECEIPTNO,0,3)"=SUBSTR(TO_CHAR("RECEIPTNO"),0,3))
   4 - filter(COUNT(*)>=10)

20 rows selected.
Elapsed: 00:00:00.04
```

For the query in task 4.1 a function-based index on SUBSTR(RECEIPTNO, 0, 3) was chosen as it is part of the WHERE clause which can yield lots of performance gain.

As seen above, the query execution time improved after making the function-based index (00.04 & 00.01 before and after the index creation respectively). This is evident within the query execution plan where before BOOK_INDEX is created the query optimizer uses the 'TABLE ACCESS FULL' while after the creation of the index the query optimizer has gone ahead and chosen to do an 'INDEX FAST FULL SCAN' on the BOOK_INDEX. This means that oracle will access the index created on SUBSTR(RECEIPTNO, 0, 3) rather than having to process the SUBSTR(). Thus improving the overall performance of the query.

4.3 Department 50

```
SELECT SUM(AMOUNT) FROM PURCHASE P
WHERE SERVEDBY IN ( SELECT EMPNO FROM EMP
                    WHERE DEPTNO = 50)
AND INSTR(SERVICETYPE, 'Software' ) = 0;
```

```
SUM(AMOUNT)
-----
905355
Elapsed: 00:00:00.05
```

4.4 Department 50 query execution plan/time and creation of index

```
EXPLAIN PLAN FOR SELECT SUM(AMOUNT) FROM PURCHASE P
WHERE SERVEDBY IN ( SELECT EMPNO FROM EMP
                    WHERE DEPTNO = 50)
AND INSTR(SERVICETYPE, 'Software' ) = 0;
```

```
Explained.
Elapsed: 00:00:00.01
```

SELECT PLAN_TABLE_OUTPUT FROM TABLE (DBMS_XPLAN.DISPLAY); --Before Index

```
PLAN_TABLE_OUTPUT
Plan hash value: 2786094069

| Id | Operation | Name | Rows | Bytes | Cost (%CPU) | Time | |
|---|---|---|---|---|---|---|---|
| 0 | SELECT STATEMENT | | 1 | 30 | 25 (4) | 00:00:01 |
| 1 | SORT AGGREGATE | | 1 | 30 | | | |
| 2 | MERGE JOIN | | 42 | 1260 | 25 (4) | 00:00:01 |
| * 3 | TABLE ACCESS BY INDEX ROWID | EMP | 24 | 168 | 2 (0) | 00:00:01 |
| 4 | INDEX FULL SCAN | PK_EMPNO | 75 | | 1 (0) | 00:00:01 |
| * 5 | SORT JOIN | | 100 | 2300 | 23 (5) | 00:00:01 |

PLAN_TABLE_OUTPUT
| * 6 | TABLE ACCESS FULL | PURCHASE | 100 | 2300 | 22 (0) | 00:00:01 |

Predicate Information (identified by operation id):
-----
 3 - filter("E"."DEPTNO"=50)
 5 - access("P"."SERVEDBY"="E"."EMPNO")
    filter("P"."SERVEDBY"="E"."EMPNO")
 6 - filter(INSTR("SERVICETYPE",'Software')=0)

21 rows selected.
Elapsed: 00:00:00.04
```

CREATE INDEX SERVICE_INDEX ON PURCHASE INSTR(SERVICETYPE, 'Software');

```
Index created.
Elapsed: 00:00:00.03
```

Result of Query After index creation.

```
SUM(AMOUNT)
-----
905355
Elapsed: 00:00:00.01
```

SELECT PLAN_TABLE_OUTPUT FROM TABLE (DBMS_XPLAN.DISPLAY); --After Index

```

PLAN_TABLE_OUTPUT
-----
Plan hash value: 3589182047

| Id | Operation                                | Name           | Rows  | Bytes | Cost (%CPU)| Time     |
|----|-----|-----|-----|-----|-----|-----|-----|
|  0 | SELECT STATEMENT                        |                |      1 |    28 |    13 (8)   | 00:00:01 |
|  1 |   SORT AGGREGATE                        |                |      1 |    28 |             |           |
|  2 |    MERGE JOIN                           |                |     42 |   1176 |    13 (8)   | 00:00:01 |
|*  3 |     TABLE ACCESS BY INDEX ROWID        | EMP            |     24 |    168 |     2 (0)   | 00:00:01 |
|  4 |      INDEX FULL SCAN                     | PK_EMPNO       |     75 |       |     1 (0)   | 00:00:01 |
|*  5 |       SORT JOIN                          |                |     100 |   2100 |    11 (10)  | 00:00:01 |

PLAN_TABLE_OUTPUT
-----
|  6 | TABLE ACCESS BY INDEX ROWID BATCHED    | PURCHASE       |     100 |   2100 |    10 (0)   | 00:00:01 |
|*  7 |   INDEX RANGE SCAN                       | SERVICE_INDEX  |      40 |       |     9 (0)   | 00:00:01 |

Predicate Information (identified by operation id):
-----
   3 - filter("E"."DEPTNO"=50)
   5 - access("P"."SERVEDBY"="E"."EMPNO")
      filter("P"."SERVEDBY"="E"."EMPNO")
   7 - access(INSTR("SERVICETYPE",'Software')=0)

22 rows selected.

Elapsed: 00:00:00.03

```

Similarly to Task 4.2 a function-based index was chosen on the INSTR(SERVICETYPE, 'Software') as it is part of the WHERE clause which can give us lots of performance gain.

Although there isn't much of a difference in the elapsed time of the query execution after the index creation (00.05 vs 00.01 before & after respectively) it is still an improvement. This is further highlighted in the query execution plan. Looking at the query execution of both before and after the index creation, the overall cost before is 98 and after is 59. This is due to operation id 7 where the oracle query optimizer uses the SERVICE_INDEX to do an 'INDEX RANGE SCAN' rather than processing the INSTR(...). Hence, improving the overall performance of the query.

4.5 ServiceType, PaymentType & GST

```

SELECT COUNT(*)
FROM PURCHASE
WHERE SERVICETYPE IN (
    SELECT SERVICETYPE FROM PURCHASE
    GROUP BY SERVICETYPE, PAYMENTTYPE, GST
    HAVING COUNT(*) >= 1000)
AND PAYMENTTYPE IN (
    SELECT PAYMENTTYPE FROM PURCHASE
    GROUP BY SERVICETYPE, PAYMENTTYPE, GST
    HAVING COUNT(*) >= 1000)
AND GST IN (
    SELECT GST FROM PURCHASE
    GROUP BY SERVICETYPE, PAYMENTTYPE, GST
    HAVING COUNT(*) >= 1000);

```

```

COUNT(*)
-----
2202
Elapsed: 00:00:00.13

```

4.6 Indexing on ServiceType, PaymentType & GST

Firstly, unlike Task 4.1~4.4 there are no function-based index that can be created within the query written in Task 4.5 thus we must choose a different type of index. Overall the most suitable type of index to be created on all three of these columns are a bitmap index. This is because there's a low number of distinct values that each column/attribute can take. SERVICETYPE, PAYMENTTYPE & GST have 5, 3 & 2 distinct values respectively as defined in the constraints earlier in Task 1.2. Additionally, DB optimizers such as the one used in Oracle, can combine several bitmap indexes very easily which will allow for efficient execution of complex filters and queries *(Another thing to mention is that although there were constraints defined on these columns to prevent null values, the bitmap index also includes null values unlike the b-tree index, thus allowing the DB optimizer to use the index on 'null' queries)*. Hence, due to these reasons the bitmap indexes will be more efficient compared to other indexes such as the b-tree index.

Task 5 – Execution Plan

5.1 PURCHASENO = 1234

```
SELECT * FROM PURCHASE WHERE PURCHASENO = 1234; --Before
```

PURCHASENO	RECEIPTNO	SERVICETYPE	PAYMENTTYPE	GST	AMOUNT	SERVEDBY	CLIENTNO
1234	591623	Consultation	Cash	Yes	860	1036	24797

Elapsed: 00:00:00.01

```
EXPLAIN PLAN FOR SELECT * FROM PURCHASE WHERE PURCHASENO = 1234;
```

```
Explained.  
Elapsed: 00:00:00.01
```

```
SELECT PLAN_TABLE_OUTPUT FROM TABLE (DBMS_XPLAN.DISPLAY); -- Before
```

```
PLAN_TABLE_OUTPUT  
-----  
Plan hash value: 2822030489  
  
| Id | Operation | Name | Rows | Bytes | Cost (%CPU)| Time |  
-----  
| 0 | SELECT STATEMENT | | 1 | 48 | 2 (0)| 00:00:01 |  
| 1 | TABLE ACCESS BY INDEX ROWID | PURCHASE | 1 | 48 | 2 (0)| 00:00:01 |  
|* 2 | INDEX UNIQUE SCAN | PK_PURCHASENO | 1 | | 1 (0)| 00:00:01 |  
-----  
  
Predicate Information (identified by operation id):  
PLAN_TABLE_OUTPUT  
-----  
  
2 - access("PURCHASENO"=1234)  
14 rows selected.  
Elapsed: 00:00:00.08
```

Execution Plan Explanation:

1. As 'INDEX UNIQUE SCAN' is the most indented to the right, this process is run first. This operation uses the index PK_PURCHASENO to search for the value within the WHERE clause which in this case is '1234' and returns just 1 row-id from the index.
2. However, all the information needed isn't in the index. Thus, the oracle kernel goes through the actual table with the pointer it got from searching through the index to fetch further information requested in the query.
3. Finally after Oracle has found the correct row, the 'SELECT STATEMENT' operation returns the row.

5.2 DROP CONSTRAINT AND RE-EXECUTE

```
ALTER TABLE PURCHASE DROP CONSTRAINT PK_PURCHASENO;
```

```
Table altered.  
Elapsed: 00:00:00.07
```

```
SELECT * FROM PURCHASE WHERE PURCHASENO = 1234; -- After
```

PURCHASENO	RECEIPTNO	SERVICETYPE	PAYMENTTYP	GST	AMOUNT	SERVEDBY	CLIENTNO
1234	591623	Consultation	Cash	Yes	860	1036	24797

```
Elapsed: 00:00:00.01
```

```
SELECT PLAN_TABLE_OUTPUT FROM TABLE (DBMS_XPLAN.DISPLAY); -- After
```

```
PLAN_TABLE_OUTPUT  
-----  
Plan hash value: 2913724801  
  
-----  
| Id | Operation          | Name      | Rows | Bytes | Cost (%CPU)| Time      |  
-----  
|  0 | SELECT STATEMENT   |           |     1 |    48 |       22 (0)| 00:00:01 |  
|*  1 |  TABLE ACCESS FULL| PURCHASE  |     1 |    48 |       22 (0)| 00:00:01 |  
-----  
  
Predicate Information (identified by operation id):  
-----  
  
PLAN_TABLE_OUTPUT  
-----  
  
1 - filter("PURCHASENO"=1234)  
13 rows selected.  
Elapsed: 00:00:00.04
```

1. This time the most indented operation is 'TABLE ACCESS FULL'. This means that every row within the table PURCHASE has been accessed to do an equality search for PURCHASENO = 1234.
2. Then the 'SELECT STATEMENT' operation returns the correct row.

Although there aren't any differences within the query elapsed time, there are differences within the execution plan that Oracle has chosen to evaluate the query. In general the main difference between the execution plan in 5.1 and 5.2 is that 5.2 uses the index thus accessing a B-tree index to search for the correct value unlike the execution plan in 5.1 which accesses the whole table and searches for the correct value by brute force. Another thing to mention is the Cost (%CPU). It shows that the overall cost for the query in 5.1 (44 overall) is much more expensive than the overall cost for the query in 5.2 (5 overall).